

Assignment 4

Rubel Roy, 706

2026-02-19

Problem to demonstrate the utility of nonlinear regression over linear regression.

Get the fgl data set from “MASS” library.

```
df1 = fgl
head(df1)

##      RI     Na    Mg    Al    Si    K    Ca   Ba    Fe type
## 1  3.01 13.64 4.49 1.10 71.78 0.06 8.75  0 0.00 WinF
## 2 -0.39 13.89 3.60 1.36 72.73 0.48 7.83  0 0.00 WinF
## 3 -1.82 13.53 3.55 1.54 72.99 0.39 7.78  0 0.00 WinF
## 4 -0.34 13.21 3.69 1.29 72.61 0.57 8.22  0 0.00 WinF
## 5 -0.58 13.27 3.62 1.24 73.08 0.55 8.07  0 0.00 WinF
## 6 -2.04 12.79 3.61 1.62 72.97 0.64 8.07  0 0.26 WinF
```

(a) Considering the refractive index (RI) of “Vehicle Window glass” as the variable of interest and assuming linearity of regression, run multiple linear regression of RI on different metallic oxides. From the p value, report which metallic oxide best explains the refractive index.

```
df1_subset = subset(df1, type == "Veh")
model = lm(RI ~ Na + Mg + Al + Si + K + Ca + Ba + Fe, data = df1_subset)
summary(model)

##
## Call:
## lm(formula = RI ~ Na + Mg + Al + Si + K + Ca + Ba + Fe, data = df1_subset)
##
## Residuals:
##      Min        1Q    Median        3Q       Max 
## -0.29194 -0.08582  0.00072  0.10740  0.33524 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 131.4641   47.2669   2.781  0.02388 *  
## Na          -0.4333    0.3509  -1.235  0.25190    
## Mg          -0.2866    1.0075  -0.285  0.78325    
## Al          -0.8909    0.5550  -1.605  0.14713    
## Si         -1.8824    0.4993  -3.770  0.00547 ** 
## K           -2.4232    0.9725  -2.492  0.03743 *  
## Ca          1.5326    0.5818   2.634  0.02998 *  
## Ba          0.3517    2.6904   0.131  0.89922    
## Fe          3.8931    0.9581   4.063  0.00362 **
```

```

## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.2621 on 8 degrees of freedom
## Multiple R-squared:  0.9906, Adjusted R-squared:  0.9813
## F-statistic: 105.9 on 8 and 8 DF,  p-value: 2.622e-07

summary(model)$coefficients

##           Estimate Std. Error   t value   Pr(>|t|) 
## (Intercept) 131.4640676 47.2669236  2.7813121 0.023876172
## Na          -0.4333080  0.3508773 -1.2349274 0.251895370
## Mg          -0.2866243  1.0074637 -0.2845009 0.783251988
## Al          -0.8908690  0.5550086 -1.6051446 0.147129402
## Si          -1.8823864  0.4993058 -3.7700067 0.005465591
## K           -2.4231984  0.9725295 -2.4916451 0.037426154
## Ca          1.5326244  0.5817872  2.6343387 0.029975590
## Ba          0.3517015  2.6904136  0.1307240 0.899221141
## Fe          3.8931318  0.9580806  4.0634699 0.003616000

stargazer(model, type = "text")

```

```

##
## =====
##             Dependent variable:
## -----
##                   RI
## -----
## Na          -0.433
##             (0.351)
## 
## Mg          -0.287
##             (1.007)
## 
## Al          -0.891
##             (0.555)
## 
## Si          -1.882***
##             (0.499)
## 
## K           -2.423**
##             (0.973)
## 
## Ca          1.533**
##             (0.582)
## 
## Ba          0.352
##             (2.690)

```

```

## 
##   Fe          3.893*** 
##                   (0.958)
## 
##   Constant      131.464** 
##                   (47.267)
## 
## -----
##   Observations       17
##   R2              0.991
##   Adjusted R2       0.981
##   Residual Std. Error    0.262 (df = 8)
##   F Statistic      105.887*** (df = 8; 8)
## -----
## Note:           *p<0.1; **p<0.05; ***p<0.01

```

Conclusion : From the p-values, it is clear that Iron (Fe) oxide has the lowest p-value and hence best explains the refractive index(RI).

(b) Run a simple linear regression of RI on the best predictor chosen in (a).

```

fit_simple = lm(RI ~ Fe, data = df1_subset)
summary(fit_simple)

## 
## Call:
## lm(formula = RI ~ Fe, data = df1_subset)
## 
## Residuals:
##     Min      1Q  Median      3Q     Max 
## -2.2324 -1.0693 -0.2715  0.2907  3.7707 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -0.5007    0.4861  -1.030   0.3193    
## Fe          8.1362    4.0780   1.995   0.0645 .  
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 1.759 on 15 degrees of freedom
## Multiple R-squared:  0.2097, Adjusted R-squared:  0.157 
## F-statistic: 3.981 on 1 and 15 DF,  p-value: 0.06452

stargazer(fit_simple, type = "text")

## -----
## Dependent variable:
## -----
##          RI

```

```

## -----
## Fe                      8.136*
##                               (4.078)
##
## Constant                 -0.501
##                               (0.486)
##
## -----
## Observations              17
## R2                       0.210
## Adjusted R2               0.157
## Residual Std. Error      1.759 (df = 15)
## F Statistic                3.981* (df = 1; 15)
## -----
## Note:                     *p<0.1; **p<0.05; ***p<0.01

```

(c) Can you further improve the regression of the refractive index of “Vehicle Window glass” on the predictor chosen by you in part (a)? Give the new fitted model and compare its performance with the model in (b).

```

fit_quad = lm(RI ~ Fe + I(Fe^2), data = df1_subset) #adding a quadratic term

summary(fit_quad)

##
## Call:
## lm(formula = RI ~ Fe + I(Fe^2), data = df1_subset)
##
## Residuals:
##     Min      1Q  Median      3Q     Max 
## -1.6215 -1.1715 -0.1345  0.5985  3.5485 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -0.2785    0.4712  -0.591   0.564    
## Fe          -12.1810   12.0408  -1.012   0.329    
## I(Fe^2)      65.9600   37.0798   1.779   0.097 .  
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.645 on 14 degrees of freedom
## Multiple R-squared:  0.3554, Adjusted R-squared:  0.2633 
## F-statistic:  3.86 on 2 and 14 DF,  p-value: 0.04623

```

Comparison :

```

cat("R^2 value for simple model : ",summary(fit_simple)$r.squared)

## R^2 value for simple model :  0.2097192

```

```

cat("R^2 value for quadratic model : ",summary(fit_quad)$r.squared)
## R^2 value for quadratic model :  0.355413

```

Conclusion : The quadratic model provides a better fit than the linear model, as shown by a higher R^2 and lower residual standard error. The overall F-test indicates statistical significance at the 5% level. However, the individual coefficients are not strongly significant, suggesting only limited evidence of a nonlinear relationship between RI and Fe.

Problem to demonstrate multicollinearity.

Consider the Credit data in the ISLR library. Choose balance as the response and Age, Limit and Rating as the predictors.

```

df2 = Credit
head(df2)

##   ID Income Limit Rating Cards Age Education Gender Student Married
Ethnicity
## 1  1 14.891  3606    283     2  34        11   Male      No    Yes
Caucasian
## 2  2 106.025  6645    483     3  82        15 Female     Yes    Yes
Asian
## 3  3 104.593  7075    514     4  71        11   Male      No    No
Asian
## 4  4 148.924  9504    681     3  36        11 Female     No    No
Asian
## 5  5 55.882   4897    357     2  68        16   Male      No    Yes
Caucasian
## 6  6 80.180   8047    569     4  77        10   Male      No    No
Caucasian
##   Balance
## 1      333
## 2      903
## 3      580
## 4      964
## 5      331
## 6     1151

```

(a) Make a scatter plot of (i) Age versus Limit and (ii) Rating Versus Limit. Comment on the scatter plot.


```

par(mfrow = c(1, 2))

plot(df2$Age, df2$Limit, main = "Age vs Limit", xlab = "Age", ylab = "Credit
Limit")

plot(df2$Rating, df2$Limit, main = "Rating vs Limit", xlab = "Rating", ylab =
"Credit Limit")

```



```
par(mfrow = c(1, 1))
```

Comments : i. *Age vs Limit*

- Type of relationship: No clear linear relationship
- Direction: No clear direction (neither upward nor downward trend)
- Strength: Very weak

The points are widely scattered without any visible pattern. This suggests that Age does not strongly explain Credit Limit.

(ii) *Rating vs Limit*

- Type of relationship: Strong linear relationship
- Direction: Positive
- Strength: Very strong

The points lie almost perfectly along a straight upward-sloping line. This indicates that as Rating increases, Limit increases almost proportionally.

(b) Run three separate regressions: (i) Balance on Age and Limit (ii) Balance on Age, Rating and Limit (iii) Balance on Rating and Limit. Present all the regression output in a single table using stargazer. What is the marked difference that you can observe from the output?

```

model21 <- lm(Balance ~ Age + Limit, data = df2)

model22 <- lm(Balance ~ Age + Rating + Limit, data = df2)

model23 <- lm(Balance ~ Rating + Limit, data = df2)

stargazer(model21, model22, model23, type = 'text', title = "Regression
Results: Balance as Response", column.labels = c("Age + Limit", "Age + Rating
+ Limit", "Rating + Limit"))

## 
## Regression Results: Balance as Response
## 
=====

##                                     Dependent variable:
##                               -----
##                               Balance
##                               Age + Limit      Age + Rating + Limit
Rating + Limit
##                               (1)          (2)
## (3)
## ----- 

## Age           -2.291***       -2.346*** 
##                   (0.672)        (0.669)
## 
## Rating        2.310** 
2.202** 
##                   (0.940)
## (0.952)
## 
## Limit         0.019
##                   0.173*** 
##                   (0.005)        (0.063)
## (0.064)
## 
## Constant     -259.518*** 
##                   -173.411*** 
##                   (43.828)        (55.882)
## (45.254)
## ----- 

```

```

## Observations          400          400
400
## R2                  0.750        0.754
0.746
## Adjusted R2         0.749        0.752
0.745
## Residual Std. Error 230.532 (df = 397)    229.080 (df = 396)
232.320 (df = 397)
## F Statistic          594.988*** (df = 2; 397) 403.718*** (df = 3; 396)
582.820*** (df = 2; 397)
##
===== =====
## Note:                      *p<0.1;
**p<0.05; ***p<0.01

```

Observations : The key difference:

1. In model (i), Limit is highly significant.
2. In model (iii), Rating is highly significant.
3. But in model (ii) (**both included**) :
 - One of them often becomes insignificant.
 - Standard errors increase.
 - Coefficient signs or magnitudes may shift.

This happens because Rating and Limit are highly correlated. That is multicollinearity.

(c) Calculate the variance inflation factor (VIF) and comment on multicollinearity.

```

vif_values <- vif(model22)
print(vif_values)

##       Age      Rating      Limit
## 1.011385 160.668301 160.592880

```

Comment on multicolinearity : There is severe multicollinearity between Rating and Limit, as indicated by extremely high VIF values (≈ 160). Age shows no multicollinearity. The instability in the regression coefficients when both Rating and Limit are included is due to this strong linear dependence.

Problem to demonstrate the detection of outlier, leverage and influential points.

Attach “Boston” data from MASS library in R. Select median value of owneroccupied homes, as the response and per capita crime rate, nitrogen oxides concentration, proportion of blacks and percentage of lower status of the population as predictors. The objective is to fit a multiple linear regression model of the response on the predictors. With reference to this problem, detect outliers, leverage points and influential points if any.

```
df3 = Boston
df3_subset = df3[, c("medv", "crim", "nox", "black", "lstat")]
head(df3_subset)

##   medv    crim    nox    black   lstat
## 1 24.0 0.00632 0.538 396.90  4.98
## 2 21.6 0.02731 0.469 396.90  9.14
## 3 34.7 0.02729 0.469 392.83  4.03
## 4 33.4 0.03237 0.458 394.63  2.94
## 5 36.2 0.06905 0.458 396.90  5.33
## 6 28.7 0.02985 0.458 394.12  5.21

boston_model = lm(medv ~ crim + nox + black + lstat, data = df3_subset)
summary(boston_model)

##
## Call:
## lm(formula = medv ~ crim + nox + black + lstat, data = df3_subset)
##
## Residuals:
##       Min     1Q   Median     3Q    Max 
## -15.564 -4.004 -1.504  2.178 24.608 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 30.053584   2.170839 13.844 <2e-16 ***
## crim        -0.059424   0.037755 -1.574   0.116    
## nox          3.415809   3.056602  1.118   0.264    
## black         0.006785   0.003408  1.991   0.047 *  
## lstat        -0.918431   0.050167 -18.307 <2e-16 ***
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.183 on 501 degrees of freedom
## Multiple R-squared:  0.5517, Adjusted R-squared:  0.5481 
## F-statistic: 154.1 on 4 and 501 DF,  p-value: < 2.2e-16
```

Residual Plot:

```
std_res = rstandard(boston_model)
```


```

plot(boston_model$fitted.values, std_res, xlab = "Fitted Values", ylab =
"Standardized Residuals", main = "Residual Plot (Standardized Residuals)")

abline(h = c(-2.5, 2.5), col = "blue", lty = 2)

```

Residual Plot (Standardized Residuals)



Calculation for Outliers :

```

outliers <- which(abs(std_res) > 2.5)

data.frame(Observation = outliers, Standardized_Residual = std_res[outliers])

##      Observation Standardized_Residual
## 162          162            2.768907
## 163          163            2.783878
## 164          164            2.988259
## 167          167            3.069974
## 187          187            3.194723
## 196          196            3.011109
## 204          204            2.897489
## 205          205            3.005419
## 215          215            2.749729
## 226          226            3.228933
## 229          229            2.591706
## 234          234            2.852734
## 258          258            3.224714
## 263          263            3.148286
## 268          268            3.605022

```

```

## 284      284      3.050345
## 369      369      3.015299
## 370      370      3.092472
## 371      371      2.970057
## 372      372      3.992694
## 373      373      3.890059
## 413      413      3.029650
## 506      506      -2.522613

```

Consistency Check : The graphical method (residual plot) and the numerical method ($|z\text{standardized residual}| > 2.5$) identify the same observations as outliers. Hence, the two approaches are consistent.

Leverage Points :

```

n = nrow(Boston)
p = length(coef(boston_model))    # includes intercept

H = hatvalues(boston_model)

threshold_lev = 2*p/n
threshold_lev

## [1] 0.01976285

leverage_points = which(H > threshold_lev)

leverage_points

##   9   33   49   103   142   143   144   145   146   147   148   149   150   151   152   153   154   155
156   157
##   9   33   49   103   142   143   144   145   146   147   148   149   150   151   152   153   154   155
156   157
## 160   215   374   375   381   386   387   388   399   401   405   406   411   412   413   414   415   416
417   418
## 160   215   374   375   381   386   387   388   399   401   405   406   411   412   413   414   415   416
417   418
## 419   420   424   425   426   427   428   430   431   432   433   434   435   437   438   439   446   451
455   456
## 419   420   424   425   426   427   428   430   431   432   433   434   435   437   438   439   446   451
455   456
## 457   458   467   491
## 457   458   467   491

```


Plot:

```

plot(H, type = "h", main = "Leverage Values (Hat Matrix Diagonal)", ylab =
"Hi")
abline(h = threshold_lev, col = "red", lty = 2)

```

Leverage Values (Hat Matrix Diagonal)



Comment : The leverage plot shows that a small number of observations exceed the leverage threshold, indicating the presence of high leverage points. Most observations have low leverage. These high leverage points may potentially influence the regression model.

Influential Points (Cook's Distance) :

```
cooks_d = cooks.distance(model)

threshold_cook = 4/n
threshold_cook

## [1] 0.007905138

influential_points = which(cooks_d > threshold_cook)

influential_points


## 147 149 150 151 152 153 155 156 157 160 161 163
##    1    3    4    5    6    7    9   10   11   14   15   17
```

Plot :

```
plot(cooks_d, type = "h", main = "Cook's Distance", ylab = "Cook's D")

abline(h = threshold_cook, col = "red", lty = 2)
```

Cook's Distance



Comment : The Cook's distance plot indicates the presence of a few influential observations, with one observation having particularly high influence. Most observations have negligible influence. The highly influential points may substantially affect the regression coefficients and should be examined further.

Effect of Removing Influential Points :

```
model_reduced = lm(medv ~ crim + nox + black + lstat, data = Boston[-influential_points, ])

summary(model_reduced)

##
## Call:
## lm(formula = medv ~ crim + nox + black + lstat, data = Boston[-influential_points,
##   ])
##
## Residuals:
##     Min      1Q      Median      3Q      Max 
## -15.645  -3.985  -1.529   2.173  24.498 
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 29.916753  2.181633 13.713 <2e-16 ***
## crim        -0.057530  0.037984 -1.515  0.1305    
## nox         4.018540  3.083138  1.303  0.1931    
```

```
## black      0.006753  0.003424  1.972   0.0491 *
## lstat     -0.933039  0.051488 -18.121  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 6.206 on 489 degrees of freedom
## Multiple R-squared:  0.5525, Adjusted R-squared:  0.5489
## F-statistic:  151 on 4 and 489 DF,  p-value: < 2.2e-16
```

Comparison to original model : After removing influential observations, the regression coefficients and overall model fit remain broadly similar. Although some minor changes in magnitude and significance occur, the general conclusions remain unchanged. Therefore, the regression model is reasonably robust to influential observations.