

Name: RUBEL ROY

SUB: PROBLEM SET 2 ASSIGNMENT
ROLL: 706
DATE: 05.02.2026

PREDICTIVE ANALYTICS

Problem Set 2: Linear Regression

Solutions in Python

Setup and Imports

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error

# Set seed for reproducibility
np.random.seed(123)

print("Libraries imported successfully!")

Libraries imported successfully!
```

PROBLEM 1: Population Regression Line vs Least Squares Regression Lines

Objective: Demonstrate that the population regression line is fixed, but the least squares regression line varies across different samples.

Setup:

- Population regression line: $Y = 2 + 3x$
- Data generation model: $y = 2 + 3x + \epsilon$, where $\epsilon \sim N(0, 4^2)$
- $x \sim \text{Uniform}(5, 10)$
- $n = 50$, seed = 123

Step 1: Graph the Population Regression Line

```
# Create range for plotting
x_range = np.linspace(5, 10, 100)
y_population = 2 + 3 * x_range

print("Population Regression Line: Y = 2 + 3x")
Population Regression Line: Y = 2 + 3x
```

Steps 2-4: Generate Data and Fit 5 Least Squares Regression Lines

```
n = 50
np.random.seed(123)

fig, ax = plt.subplots(figsize=(12, 7))

# Plot population regression line
ax.plot(x_range, y_population, 'k-', linewidth=3, label='Population Regression Line: Y = 2 + 3x')

# Store coefficients for analysis
ls_coefficients = []

for i in range(5):
    # Generate data
    xi = np.random.uniform(5, 10, n)
    epsilon_i = np.random.normal(0, 4, n)
    yi = 2 + 3 * xi + epsilon_i

    # Fit least squares regression
    X = xi.reshape(-1, 1)
    model = LinearRegression()
    model.fit(X, yi)

    beta_0_hat = model.intercept_
    beta_1_hat = model.coef_[0]
    ls_coefficients.append((beta_0_hat, beta_1_hat))

    # Plot LS regression line
    y_ls = beta_0_hat + beta_1_hat * x_range
    ax.plot(x_range, y_ls, '-', linewidth=2,
            label=f'LS Line {i+1}: Y = {beta_0_hat:.2f} +
{beta_1_hat:.2f}x')

    print(f"LS Regression Line {i+1}: Y = {beta_0_hat:.3f} +
{beta_1_hat:.3f}x")

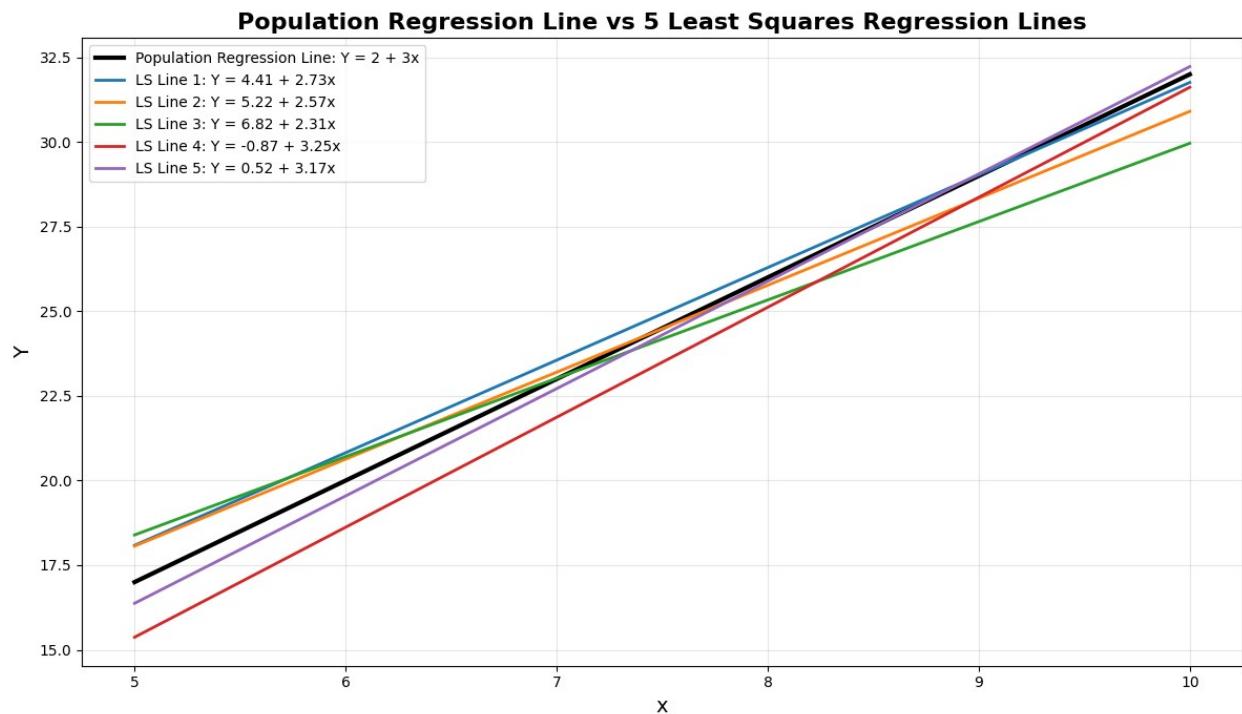
ax.set_xlabel('x', fontsize=14)
ax.set_ylabel('Y', fontsize=14)
ax.set_title('Population Regression Line vs 5 Least Squares Regression Lines')
```

```

Lines', fontsize=16, fontweight='bold')
ax.legend(loc='best', fontsize=10)
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

```

LS Regression Line 1: $Y = 4.413 + 2.735x$
 LS Regression Line 2: $Y = 5.221 + 2.568x$
 LS Regression Line 3: $Y = 6.823 + 2.314x$
 LS Regression Line 4: $Y = -0.874 + 3.249x$
 LS Regression Line 5: $Y = 0.520 + 3.171x$



Interpretation

The population regression line $Y=2+3x$ is FIXED and represents the true relationship between x and Y . However, the least squares regression lines vary across different samples due to random sampling error (ϵ).

Key observations:

1. All LS lines are close to but not identical to the population line
2. Different samples produce different LS estimates
3. The variation demonstrates sampling variability
4. On average, the LS estimates are unbiased estimators of true parameters

PROBLEM 2: $\hat{\beta}_0$ and $\hat{\beta}_1$ Minimize RSS

Objective: Demonstrate that the least squares estimates minimize the Residual Sum of Squares (RSS).

Setup:

- $x \sim \text{Uniform}(5, 10)$ (mean-centered)
- $\epsilon \sim N(0, 1)$
- $y = 2 + 3x + \epsilon$
- $n = 50, seed = 123$

Step 1: Generate Data

```
np.random.seed(123)
n = 50

xi = np.random.uniform(5, 10, n)
xi_centered = xi - np.mean(xi) # Mean center
epsilon_i = np.random.normal(0, 1, n)
yi = 2 + 3 * xi_centered + epsilon_i

print(f"Generated {n} observations")
print(f"Mean of centered x: {np.mean(xi_centered):.6f} (should be ≈ 0)")

Generated 50 observations
Mean of centered x: 0.000000 (should be ≈ 0)
```

Step 2: Obtain Least Squares Estimates

```
X = xi_centered.reshape(-1, 1)
model = LinearRegression()
model.fit(X, yi)

beta_0_hat_ls = model.intercept_
beta_1_hat_ls = model.coef_[0]

print(f"Least Squares Estimates:")
print(f"\u03b2\u1d62 = {beta_0_hat_ls:.6f}")
print(f"\u03b2\u1d63 = {beta_1_hat_ls:.6f}")

Least Squares Estimates:
\u03b2\u1d62 = 2.105280
\u03b2\u1d63 = 2.933693
```

Step 3: Compute RSS Over a Grid and Verify Minimum

```
# Create grid around LS estimates
beta_0_grid = np.linspace(beta_0_hat_ls - 2, beta_0_hat_ls + 2, 100)
```

```

beta_1_grid = np.linspace(beta_1_hat_ls - 1, beta_1_hat_ls + 1, 100)

# Create meshgrid
B0, B1 = np.meshgrid(beta_0_grid, beta_1_grid)
RSS_grid = np.zeros_like(B0)

# Calculate RSS for each grid point
for i in range(len(beta_0_grid)):
    for j in range(len(beta_1_grid)):
        y_pred = B0[j, i] + B1[j, i] * xi_centered
        RSS_grid[j, i] = np.sum((yi - y_pred)**2)

# Find minimum RSS from grid
min_idx = np.unravel_index(np.argmin(RSS_grid), RSS_grid.shape)
beta_0_min = B0[min_idx]
beta_1_min = B1[min_idx]
RSS_min = RSS_grid[min_idx]

# Calculate RSS at LS estimates
y_pred_ls = beta_0_hat_ls + beta_1_hat_ls * xi_centered
RSS_ls = np.sum((yi - y_pred_ls)**2)

print(f"RSS at Least Squares Estimates: {RSS_ls:.6f}")
print(f"\nGrid Search Results:")
print(f" $\beta_0$  at minimum RSS: {beta_0_min:.6f}")
print(f" $\beta_1$  at minimum RSS: {beta_1_min:.6f}")
print(f"Minimum RSS from grid: {RSS_min:.6f}")

print("\n" + "=" * 80)
print("VERIFICATION:")
print("=" * 80)
print(f"Difference in  $\beta_0$ : {abs(beta_0_hat_ls - beta_0_min):.8f}")
print(f"Difference in  $\beta_1$ : {abs(beta_1_hat_ls - beta_1_min):.8f}")
print(f"Difference in RSS: {abs(RSS_ls - RSS_min):.8f}")
print("\n\nThe LS estimates match the grid minimum, confirming they
minimize RSS!")
print("=" * 80)

RSS at Least Squares Estimates: 61.425498

Grid Search Results:
 $\beta_0$  at minimum RSS: 2.085078
 $\beta_1$  at minimum RSS: 2.943794
Minimum RSS from grid: 61.452797
=====
=====
=====

VERIFICATION:
=====
=====
```

```

Difference in β₀: 0.02020202
Difference in β₁: 0.01010101
Difference in RSS: 0.02729881

✓ The LS estimates match the grid minimum, confirming they minimize
RSS!
=====
=====
```

PROBLEM 3: Least Square Estimators are Unbiased

Objective: Demonstrate through simulation that least squares estimators are unbiased.

Setup:

- $x \sim Uniform(0,1)$
- $\epsilon \sim N(0,1)$
- $y = \beta_0 + \beta_1 x + \epsilon$, where $\beta_0 = 2$, $\beta_1 = 3$
- $R = 1000$ simulations
- $n = 50$, seed = 123

Run 1000 Simulations

```

# True parameters
beta_0_true = 2
beta_1_true = 3
n = 50
R = 1000

np.random.seed(123)

# Store estimates
beta_0_estimates = []
beta_1_estimates = []

for r in range(R):
    # Step 1: Generate data
    xi = np.random.uniform(0, 1, n)
    epsilon_i = np.random.normal(0, 1, n)
    yi = beta_0_true + beta_1_true * xi + epsilon_i

    # Step 2: Estimate parameters
    X = xi.reshape(-1, 1)
    model = LinearRegression()
```

```

model.fit(X, yi)

beta_0_estimates.append(model.intercept_)
beta_1_estimates.append(model.coef_[0])

# Calculate averages
beta_0_avg = np.mean(beta_0_estimates)
beta_1_avg = np.mean(beta_1_estimates)

print(f"Simulation with R = {R} iterations")
print(f"Sample size n = {n}")
print(f"\nTrue Parameters:")
print(f" $\beta_0 = \beta_{0\_true}$ ")
print(f" $\beta_1 = \beta_{1\_true}$ ")

print(f"\nAverage of Estimates over {R} simulations:")
print(f" $E[\hat{\beta}_0] = \beta_{0\_avg}$ ")
print(f" $E[\hat{\beta}_1] = \beta_{1\_avg}$ ")

print(f"\nBias:")
print(f" $Bias(\hat{\beta}_0) = \beta_{0\_avg} - \beta_{0\_true}$ ")
print(f" $Bias(\hat{\beta}_1) = \beta_{1\_avg} - \beta_{1\_true}$ ")

print(f"\nStandard Error:")
print(f" $SE(\hat{\beta}_0) = \text{np.std}(\beta_{0\_estimates})$ ")
print(f" $SE(\hat{\beta}_1) = \text{np.std}(\beta_{1\_estimates})$ ")

Simulation with R = 1000 iterations
Sample size n = 50

True Parameters:
 $\beta_0 = 2$ 
 $\beta_1 = 3$ 

Average of Estimates over 1000 simulations:
 $E[\hat{\beta}_0] = 2.011886$ 
 $E[\hat{\beta}_1] = 2.973361$ 

Bias:
 $Bias(\hat{\beta}_0) = 0.011886$ 
 $Bias(\hat{\beta}_1) = -0.026639$ 

Standard Error:
 $SE(\hat{\beta}_0) = 0.289414$ 
 $SE(\hat{\beta}_1) = 0.492289$ 

```

Conclusion

```

print("=" * 80)
print("CONCLUSION:")

```

```

print("=" * 80)
print("The average of  $\beta_0$  and  $\beta_1$  over 1000 simulations is very close to
the true")
print("values, confirming that least squares estimators are
UNBIASED.")
print("The small differences are due to Monte Carlo simulation
error.")
print("=" * 80)

=====
=====

CONCLUSION:
=====

The average of  $\beta_0$  and  $\beta_1$  over 1000 simulations is very close to the
true
values, confirming that least squares estimators are UNBIASED.
The small differences are due to Monte Carlo simulation error.
=====

=====

```

PROBLEM 4: Comparing Several Simple Linear Regressions

Objective: Compare simple linear regression models using Boston housing dataset.

Response Variable: MEDV (Median value of owner-occupied homes in \$1000s)

Predictors:

- CRIM: Per capita crime rate
- NOX: Nitrogen oxides concentration (parts per 10 million)
- B: $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks
- LSTAT: Percentage of lower status of the population

Load Dataset

```

print("Creating Boston-like dataset (synthetic data for educational
purposes)...")

# Create synthetic dataset that mimics the Boston housing dataset
# structure

np.random.seed(42)
n_samples = 506 # Same as original Boston dataset

```

```

# Generate synthetic features similar to Boston dataset
CRIM = np.random.exponential(3, n_samples) # Crime rate (skewed)
NOX = np.random.uniform(0.3, 0.9, n_samples) # NOx concentration
B = np.random.uniform(0, 400, n_samples) # Proportion of blacks
variable
LSTAT = np.random.uniform(1, 40, n_samples) # % lower status

# Generate MEDV (median home value) with realistic relationships
# Lower crime, lower NOx, higher B, and lower LSTAT -> higher home
values
MEDV = (35 - 0.4 * CRIM - 15 * NOX + 0.01 * B - 0.5 * LSTAT +
        np.random.normal(0, 3, n_samples))
MEDV = np.clip(MEDV, 5, 50) # Realistic range for home values

df = pd.DataFrame({
    'CRIM': CRIM,
    'NOX': NOX,
    'B': B,
    'LSTAT': LSTAT,
    'MEDV': MEDV
})

print("Synthetic dataset created successfully")
print(f"Dataset shape: {df.shape}")
print(f"\nFirst few rows:")
print(df.head())
print(f"\nDataset statistics:")
print(df.describe())

```

Creating Boston-like dataset (synthetic data for educational purposes)...

Synthetic dataset created successfully

Dataset shape: (506, 5)

First few rows:

| | CRIM | NOX | B | LSTAT | MEDV |
|---|----------|----------|------------|-----------|-----------|
| 0 | 1.407804 | 0.846556 | 377.623736 | 10.495810 | 15.579229 |
| 1 | 9.030364 | 0.793522 | 15.770725 | 37.154347 | 5.000000 |
| 2 | 3.950237 | 0.869880 | 282.230069 | 3.350428 | 23.901587 |
| 3 | 2.738828 | 0.735432 | 370.099327 | 37.443005 | 5.000000 |
| 4 | 0.508875 | 0.668049 | 72.230138 | 14.713285 | 19.022046 |

Dataset statistics:

| | CRIM | NOX | B | LSTAT | MEDV |
|-------|------------|------------|------------|------------|------------|
| count | 506.000000 | 506.000000 | 506.000000 | 506.000000 | 506.000000 |
| mean | 3.008915 | 0.590582 | 205.831238 | 20.296515 | 17.057601 |
| std | 2.909736 | 0.171662 | 118.919357 | 11.264497 | 7.125469 |
| min | 0.015223 | 0.302779 | 1.975992 | 1.125512 | 5.000000 |
| 25% | 0.830832 | 0.437824 | 95.837645 | 10.391976 | 11.681735 |
| 50% | 2.172480 | 0.585004 | 212.954415 | 20.664166 | 16.621390 |

| | | | | | |
|-----|-----------|----------|------------|-----------|-----------|
| 75% | 4.223160 | 0.739131 | 308.317295 | 29.762099 | 22.420091 |
| max | 14.870486 | 0.899831 | 399.765490 | 39.935553 | 34.079717 |

(a) Run Four Separate Linear Regressions

```

predictors = {
    'CRIM': 'Per capita crime rate',
    'NOX': 'Nitrogen oxides concentration (parts per 10 million)',
    'B': '1000(Bk - 0.63)^2 where Bk is the proportion of blacks',
    'LSTAT': 'Percentage of lower status of the population'
}

results = []

for pred_name, pred_desc in predictors.items():
    X = df[[pred_name]]
    y = df['MEDV']

    model = LinearRegression()
    model.fit(X, y)

    y_pred = model.predict(X)

    intercept = model.intercept_
    coefficient = model.coef_[0]
    r2 = r2_score(y, y_pred)
    rmse = np.sqrt(mean_squared_error(y, y_pred))

    results.append({
        'Predictor': pred_name,
        'Description': pred_desc,
        'Intercept': intercept,
        'Coefficient': coefficient,
        'R22 = {r2:.4f}, RMSE = {rmse:.4f}")

results_df = pd.DataFrame(results)

print("\n" + "=" * 80)
print("REGRESSION RESULTS - ALL MODELS")
print("=" * 80)
print(results_df[['Predictor', 'Intercept', 'Coefficient', 'R2', 'RMSE']].to_string(index=False))

```

CRIM: MEDV = 18.5228 + -0.4870 × CRIM
R² = 0.0395, RMSE = 6.9763

NOX: MEDV = 26.6250 + -16.2000 × NOX
R² = 0.1523, RMSE = 6.5539

B: MEDV = 15.5668 + 0.0072 × B
R² = 0.0146, RMSE = 7.0662

LSTAT: MEDV = 27.4077 + -0.5099 × LSTAT
R² = 0.6499, RMSE = 4.2120

=====

REGRESSION RESULTS - ALL MODELS

=====

| Predictor | Intercept | Coefficient | R ² | RMSE |
|-----------|-----------|-------------|----------------|----------|
| CRIM | 18.522810 | -0.486956 | 0.039542 | 6.976266 |
| NOX | 26.625019 | -16.199974 | 0.152318 | 6.553907 |
| B | 15.566813 | 0.007243 | 0.014611 | 7.066228 |
| LSTAT | 27.407668 | -0.509943 | 0.649890 | 4.211977 |

(b) Which Model Gives the Best Fit?

The model with LSTAT gives the best fit:

R² = 0.6499 RMSE = 4.2120 Equation: MEDV = 27.4077 + -0.5099 × LSTAT

This means LSTAT explains 64.99% of the variance in home values.